

May 20, 2024

1 Practical Question 1

Ques1. Create a file “people.txt” with the following data:

Age agegroup height status yearsmarried

21 adult 6.0 single -1

2 child 3 married 0

18 adult 5.7 married 20

221 elderly 5 widowed 2

34 child -7 married 3

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

i) Read the data from the file “people.txt”.

```
[2]: df=pd.read_csv('people.txt')
df
```

```
[2]:
```

| | Age | agegroup | height | status | yearsmarried |
|---|-----|----------|--------|---------|--------------|
| 0 | 21 | adult | 6.0 | single | -1 |
| 1 | 2 | child | 3.0 | married | 0 |
| 2 | 18 | adult | 5.7 | married | 20 |
| 3 | 221 | elderly | 5.0 | widowed | 2 |
| 4 | 34 | child | -7.0 | married | 3 |

```
[3]: df1 = np.array(df)
for i in df1:
    print(i)
```

```
[21 'adult' 6.0 'single' -1]
[2 'child' 3.0 'married' 0]
[18 'adult' 5.7 'married' 20]
[221 'elderly' 5.0 'widowed' 2]
[34 'child' -7.0 'married' 3]
```

ii) Create a ruleset E that contain rules to check for the following conditions: 1. The age should be in the range 0-150. 2. The age should be greater than yearsmarried. 3. The status should be married or single or widowed. 4. If age is less than 18 the agegroup should be child, if age is between 18 and 65 the agegroup should be adult, if age is more than 65 the agegroup should be elderly.

```
[4]: def ruleset_E(data):
    violations = []
    for record in data:
        Age, agegroup, height, status, yearsmarried = record
        if not (0 <= Age <= 150):
            violations.append(f"Age {Age} is not in the range 0-150.")
        if Age <= yearsmarried:
            violations.append(f"Age {Age} is not greater than years married_
↳({yearsmarried}).")
        if status not in ['single', 'married', 'widowed']:
            violations.append(f"Status {status} is not valid.")
        if Age < 18 and agegroup != 'child':
            violations.append(f"Age {Age} should have agegroup 'child'.")
        elif 18 <= Age <= 65 and agegroup != 'adult':
            violations.append(f"Age {Age} should have agegroup 'adult'.")
        elif Age > 65 and agegroup != 'elderly':
            violations.append(f"Age {Age} should have agegroup 'elderly'.")
    return violations
```

iii) Check whether ruleset E is violated by the data in the file people.txt.

```
[5]: violations=ruleset_E(df1)
violations
```

```
[5]: ['Age 18 is not greater than years married (20).',
      'Age 221 is not in the range 0-150.',
      "Age 34 should have agegroup 'adult'."]
```

iv) Summarize the results obtained in part (iii)

```
[6]: if violations:
    print("Ruleset E is violated by the data in people.txt:")
    for violation in violations:
        print(violation)
else:
    print("No violations found in the data.")
```

Ruleset E is violated by the data in people.txt:
Age 18 is not greater than years married (20).
Age 221 is not in the range 0-150.
Age 34 should have agegroup 'adult'.

v) Visualize the results obtained in part (iii)

```

[7]: import pandas as pd
import matplotlib.pyplot as plt

# Read data from file using read_csv method
def read_data_csv(file_path):
    df = pd.read_csv(file_path, delimiter=' ')
    return df.values.tolist()

# Define ruleset E
def ruleset_E(data):
    violations_count = {'Age range': 0, 'Age greater than years married': 0,
↳ 'Valid status': 0, 'Valid age group': 0}
    for record in data:
        age, agegroup, height, status, yearsmarried = record
        if not (0 <= age <= 150):
            violations_count['Age range'] += 1
        if age <= yearsmarried:
            violations_count['Age greater than years married'] += 1
        if status not in ['single', 'married', 'widowed']:
            violations_count['Valid status'] += 1
        if (age < 18 and agegroup != 'child') or (18 <= age <= 65 and agegroup !
↳ 'adult') or (age > 65 and agegroup != 'elderly'):
            violations_count['Valid age group'] += 1
    return violations_count

# Read data from file
file_path = "people.txt"
data = read_data_csv(file_path)

# Apply ruleset E
violations_count = ruleset_E(df1)

# Summarize results
if sum(violations_count.values()) > 0:
    print("Ruleset E is violated by the data in people.txt:")
    for rule, count in violations_count.items():
        if count > 0:
            print(f"Rule '{rule}': {count} violation(s)")
else:
    print("No violations found in the data.")

# Visualize results
rules = list(violations_count.keys())
counts = list(violations_count.values())

plt.figure(figsize=(10, 6))
plt.bar(rules, counts, color='skyblue')

```

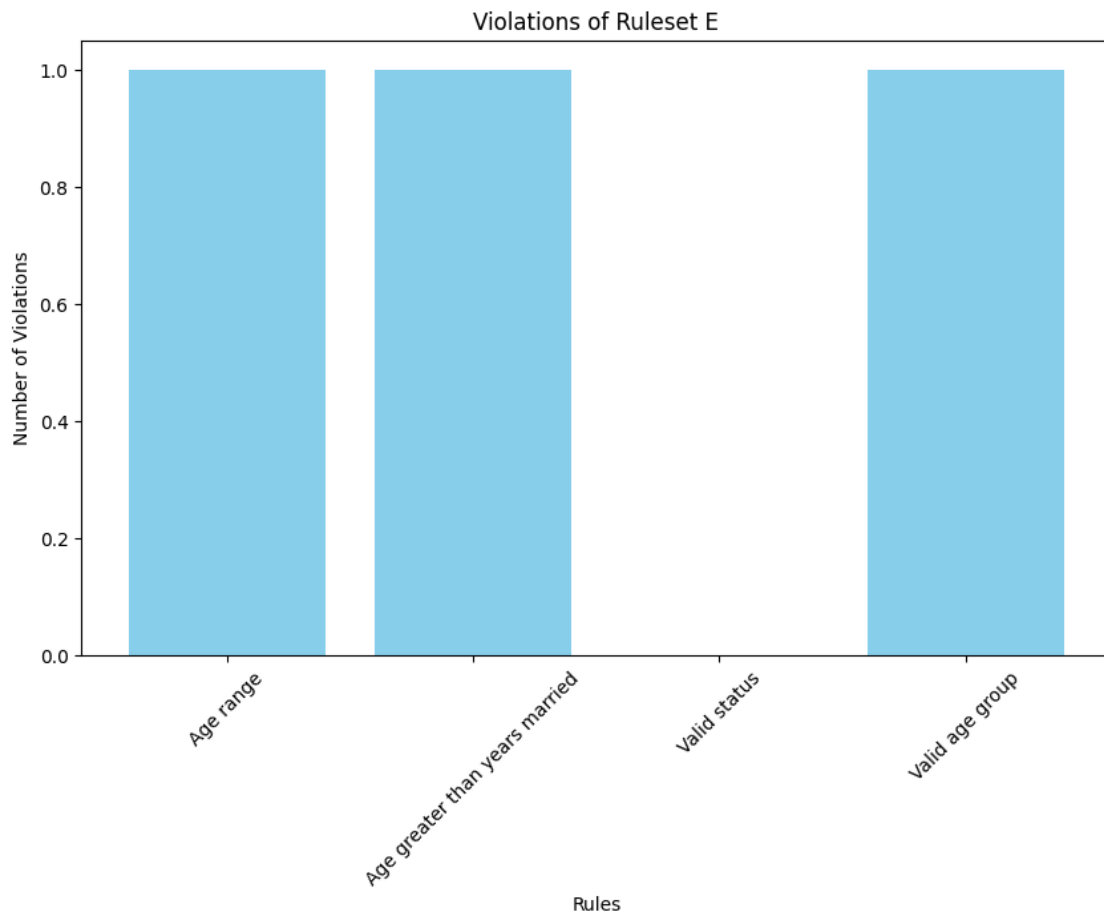
```
plt.xlabel('Rules')
plt.ylabel('Number of Violations')
plt.title('Violations of Ruleset E')
plt.xticks(rotation=45)
plt.show()
```

Ruleset E is violated by the data in people.txt:

Rule 'Age range': 1 violation(s)

Rule 'Age greater than years married': 1 violation(s)

Rule 'Valid age group': 1 violation(s)



[]:

[]:

2 Practical Question 2

Ques2. Perform the following preprocessing tasks on the dirty_iris datasetii.

```
[8]: import pandas as pd
import numpy as np
df2 = pd.DataFrame(pd.read_csv("dirty_iris.csv"))
df2
```

```
[8]:      Unnamed: 0  Sepal.Length  Sepal.Width  Petal.Length  Petal.Width  \
0              1           5.1           3.5           1.4           0.2
1              2           4.9           3.0           1.4           0.2
2              3           4.7           3.2           1.3           0.2
3              4           4.6           3.1           1.5           0.2
4              5           NaN           3.6           1.4           0.2
..          ...          ...          ...          ...          ...
145          146           6.7           3.0           5.2           2.3
146          147           6.3           2.5           5.0           1.9
147          148           6.5           3.0           5.2           2.0
148          149           6.2           3.4           5.4           2.3
149          150           NaN           3.0           5.1           1.8

      Species
0      Setosa
1      setosa
2      setosa
3      setosa
4      setosa
..          ...
145  virginica
146  virginica
147  virginica
148  virginica
149  virginica

[150 rows x 6 columns]
```

i) Calculate the number and percentage of observations that are complete.

```
[9]: df1 = df2.dropna()
print("No of rows which are free of null values : ", df1.shape[0])
print("Percentage of rows which are free of null values : ", (df1.shape[0]/df2.
↪shape[0])*100, "%")
```

No of rows which are free of null values : 131

Percentage of rows which are free of null values : 87.33333333333333 %

ii) Replace all the special values in data with NA.

```
[10]: df3=df2.replace(to_replace=np.NaN,value="NA")
df3
```

```
[10]:      Unnamed: 0 Sepal.Length Sepal.Width Petal.Length Petal.Width Species
0          1      5.1          3.5          1.4          0.2    Setosa
1          2      4.9          3.0          1.4          0.2    setosa
2          3      4.7          3.2          1.3          0.2    setosa
3          4      4.6          3.1          1.5          0.2    setosa
4          5      NA          3.6          1.4          0.2    setosa
..      ...      ...      ...      ...      ...
145      146      6.7          3.0          5.2          2.3  virginica
146      147      6.3          2.5          5.0          1.9  virginica
147      148      6.5          3.0          5.2          2.0  virginica
148      149      6.2          3.4          5.4          2.3  virginica
149      150      NA          3.0          5.1          1.8  virginica
```

[150 rows x 6 columns]

iii) Define these rules in a separate text file and read them.

(Use editfile function in R (package editrules). Use similar function in Python).

Print the resulting constraint object.

- Species should be one of the following values: setosa, versicolor or virginica.
- All measured numerical properties of an iris should be positive.
- The petal length of an iris is at least 2 times its petal width.
- The sepal length of an iris cannot exceed 30 cm.
- The sepals of an iris are longer than its petals.

```
[11]: def species():
    print("Species:")
    print((df2["Species"]=="versicolor") | (df2["Species"]=="virginica") |
    ↪(df2["Species"]=="setosa"))
    x=((df2["Species"]=="versicolor") | (df2["Species"]=="virginica") |
    ↪(df2["Species"]=="setosa"))
    print("no.of valid values:",x.sum())
    print("no.of valid values:",len(x)-x.sum())
```

```
[12]: def positive():
    print("Positive value:")
    print((df2["Sepal.Length"]>0) & (df2["Sepal.Width"]>0) & (df2["Petal.
    ↪Length"]>0) & (df2["Petal.Width"]>0))
    x=((df2["Sepal.Length"]>0) & (df2["Sepal.Width"]>0) & (df2["Petal.
    ↪Length"]>0) & (df2["Petal.Width"]>0))
    print("no.of valid values:",x.sum())
    print("no.of valid values:",len(x)-x.sum())
```

```
[13]: def petal_len():
    print("Petal length:")
    print(df2["Petal.Length"]>=2*df2["Petal.Width"])
    x=(df2["Petal.Length"]>=2*df2["Petal.Width"])
    print("no.of valid values:",x.sum())
    print("no.of valid values:",len(x)-x.sum())
```

```
[14]: def sepal_petal():
    print("Sepal petal:")
    print(df2["Sepal.Length"]>df2["Petal.Length"])
    x=(df2["Sepal.Length"]>df2["Petal.Length"])
    print("no.of valid values:",x.sum())
    print("no.of valid values:",len(x)-x.sum())
```

```
[15]: def sepal_len():
    print("sepal length:")
    print(df2["Sepal.Length"]<30)
    x=(df2["Sepal.Length"]<30)
    print("no.of valid values:",x.sum())
    print("no.of valid values:",len(x)-x.sum())
```

```
[16]: results={species(),positive(),petal_len(),sepal_petal(),sepal_petal()}
```

```
Species:
0      False
1       True
2       True
3       True
4       True
...
145     True
146     True
147     True
148     True
149     True
Name: Species, Length: 150, dtype: bool
no.of valid values: 145
no.of valid values: 5
Positive value:
0       True
1       True
2       True
3       True
4      False
...
145     True
146     True
```

```

147      True
148      True
149     False
Length: 150, dtype: bool
no.of valid values: 131
no.of valid values: 19
Petal length:
0       True
1       True
2       True
3       True
4       True
...
145     True
146     True
147     True
148     True
149     True
Length: 150, dtype: bool
no.of valid values: 150
no.of valid values: 0
Sepal petal:
0       True
1       True
2       True
3       True
4      False
...
145     True
146     True
147     True
148     True
149     False
Length: 150, dtype: bool
no.of valid values: 142
no.of valid values: 8
Sepal petal:
0       True
1       True
2       True
3       True
4      False
...
145     True
146     True
147     True
148     True
149     False

```



Length: 150, dtype: bool
no.of valid values: 142
no.of valid values: 8

iv) Determine how often each rule is broken (violatedEdits). Also summarize and plot the result.

```
[17]: # Summarize violations
violations_sum = df1.sum()
print(violations_sum)

# Plot violations
df1.plot(kind='bar')
plt.xlabel('Rule')
plt.ylabel('Number of Violations')
plt.title('Number of Violations per Rule')
plt.show()
```

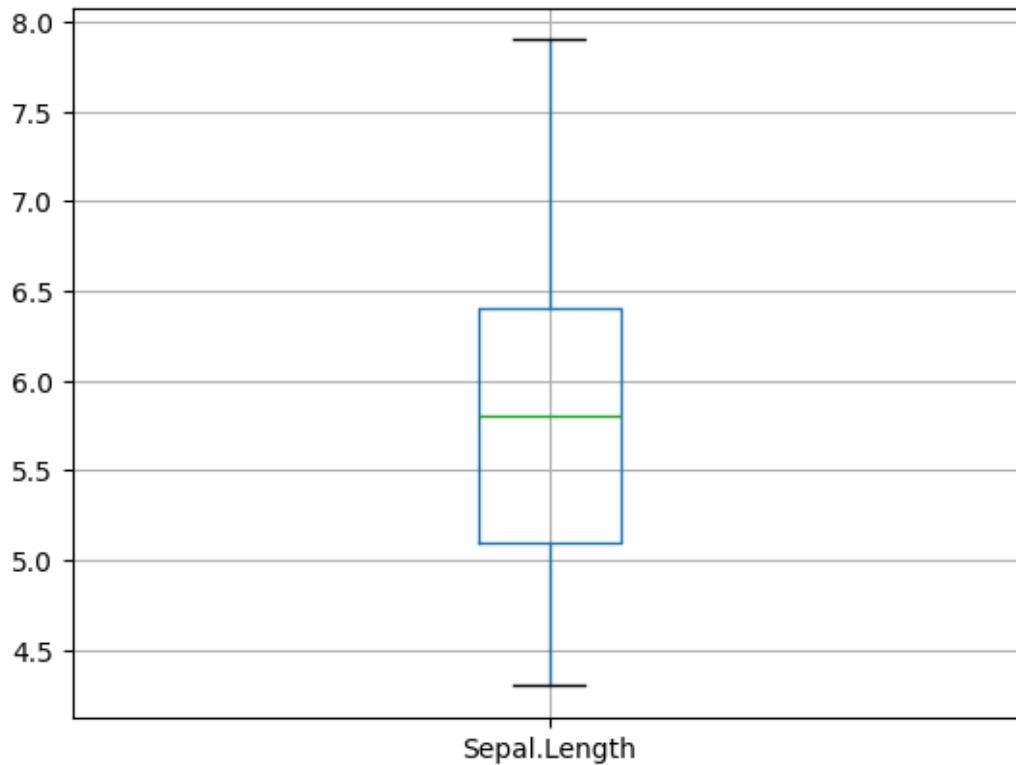
| | |
|--------------|---|
| Unnamed: 0 | 9864 |
| Sepal.Length | 766.1 |
| Sepal.Width | 397.8 |
| Petal.Length | 496.2 |
| Petal.Width | 157.6 |
| Species | SetosasetosasetosasetosasetosaSETOSasetosaseto... |
| dtype: | object |



v) Find outliers in sepal length using boxplot and boxplot.stats

```
[18]: df1[["Sepal.Length"]].boxplot()
```

```
[18]: <Axes: >
```



[]:

[]:

3 Practical Question 3

Ques3: Load the data from wine dataset. Check whether all attributes are standardized or not (mean is 0 and standard deviation is 1). If not, standardize the attributes. Do the same with Iris dataset.

```
[19]: import numpy as np
import pandas as pd
from sklearn import preprocessing

df=pd.read_csv('wine.csv')
df
```

```
[19]:
```

| | Wine | Alcohol | Malic.acid | Ash | Ac1 | Mg | Phenols | Flavanoids | \ |
|---|------|---------|------------|------|------|-----|---------|------------|---|
| 0 | 1 | 14.23 | 1.71 | 2.43 | 15.6 | 127 | 2.80 | 3.06 | |
| 1 | 1 | 13.20 | 1.78 | 2.14 | 11.2 | 100 | 2.65 | 2.76 | |
| 2 | 1 | 13.16 | 2.36 | 2.67 | 18.6 | 101 | 2.80 | 3.24 | |
| 3 | 1 | 14.37 | 1.95 | 2.50 | 16.8 | 113 | 3.85 | 3.49 | |

| | | | | | | | | |
|-----|-----|-------|------|------|------|-----|------|------|
| 4 | 1 | 13.24 | 2.59 | 2.87 | 21.0 | 118 | 2.80 | 2.69 |
| .. | ... | ... | ... | ... | ... | ... | ... | ... |
| 173 | 3 | 13.71 | 5.65 | 2.45 | 20.5 | 95 | 1.68 | 0.61 |
| 174 | 3 | 13.40 | 3.91 | 2.48 | 23.0 | 102 | 1.80 | 0.75 |
| 175 | 3 | 13.27 | 4.28 | 2.26 | 20.0 | 120 | 1.59 | 0.69 |
| 176 | 3 | 13.17 | 2.59 | 2.37 | 20.0 | 120 | 1.65 | 0.68 |
| 177 | 3 | 14.13 | 4.10 | 2.74 | 24.5 | 96 | 2.05 | 0.76 |

| | Nonflavanoid.phenols | Proanth | Color.int | Hue | OD | Proline |
|-----|----------------------|---------|-----------|------|------|---------|
| 0 | 0.28 | 2.29 | 5.64 | 1.04 | 3.92 | 1065 |
| 1 | 0.26 | 1.28 | 4.38 | 1.05 | 3.40 | 1050 |
| 2 | 0.30 | 2.81 | 5.68 | 1.03 | 3.17 | 1185 |
| 3 | 0.24 | 2.18 | 7.80 | 0.86 | 3.45 | 1480 |
| 4 | 0.39 | 1.82 | 4.32 | 1.04 | 2.93 | 735 |
| .. | ... | ... | ... | ... | ... | ... |
| 173 | 0.52 | 1.06 | 7.70 | 0.64 | 1.74 | 740 |
| 174 | 0.43 | 1.41 | 7.30 | 0.70 | 1.56 | 750 |
| 175 | 0.43 | 1.35 | 10.20 | 0.59 | 1.56 | 835 |
| 176 | 0.53 | 1.46 | 9.30 | 0.60 | 1.62 | 840 |
| 177 | 0.56 | 1.35 | 9.20 | 0.61 | 1.60 | 560 |

[178 rows x 14 columns]

```
[20]: df.std()
```

```
[20]: Wine                0.775035
      Alcohol            0.811827
      Malic.acid         1.117146
      Ash                0.274344
      Acl                3.339564
      Mg                14.282484
      Phenols            0.625851
      Flavanoids         0.998859
      Nonflavanoid.phenols 0.124453
      Proanth            0.572359
      Color.int          2.318286
      Hue                0.228572
      OD                0.709990
      Proline           314.907474
      dtype: float64
```

```
[21]: df.mean()
```

```
[21]: Wine                1.938202
      Alcohol           13.000618
      Malic.acid        2.336348
      Ash               2.366517
```

```

Ac1                19.494944
Mg                 99.741573
Phenols            2.295112
Flavanoids         2.029270
Nonflavanoid.phenols 0.361854
Proanth            1.590899
Color.int          5.058090
Hue                0.957449
OD                 2.611685
Proline            746.893258
dtype: float64

```

```

[22]: df1 = (df-df.mean())/df.std()
df1

```

```

[22]:
      Wine  Alcohol  Malic.acid  Ash  Ac1  Mg  Phenols  \
0  -1.210529  1.514341  -0.560668  0.231400 -1.166303  1.908522  0.806722
1  -1.210529  0.245597  -0.498009 -0.825667 -2.483841  0.018094  0.567048
2  -1.210529  0.196325   0.021172  1.106214 -0.267982  0.088110  0.806722
3  -1.210529  1.686791  -0.345835  0.486554 -0.806975  0.928300  2.484437
4  -1.210529  0.294868   0.227053  1.835226  0.450674  1.278379  0.806722
..      ...      ...      ...      ...      ...      ...
173  1.370000  0.873810   2.966176  0.304301  0.300954 -0.331985 -0.982841
174  1.370000  0.491955   1.408636  0.413653  1.049555  0.158126 -0.791103
175  1.370000  0.331822   1.739837 -0.388260  0.151234  1.418411 -1.126646
176  1.370000  0.208643   0.227053  0.012696  0.151234  1.418411 -1.030776
177  1.370000  1.391162   1.578712  1.361368  1.498716 -0.261969 -0.391646

      Flavanoids  Nonflavanoid.phenols  Proanth  Color.int  Hue  \
0      1.031908      -0.657708  1.221438   0.251009  0.361158
1      0.731565      -0.818411 -0.543189  -0.292496  0.404908
2      1.212114      -0.497005  2.129959   0.268263  0.317409
3      1.462399      -0.979113  1.029251   1.182732 -0.426341
4      0.661485      0.226158  0.400275  -0.318377  0.361158
..      ...      ...      ...      ...      ...
173  -1.420891      1.270726 -0.927563   1.139596 -1.388840
174  -1.280731      0.547563 -0.316058   0.967055 -1.126341
175  -1.340800      0.547563 -0.420888   2.217979 -1.607590
176  -1.350811      1.351077 -0.228701   1.829761 -1.563840
177  -1.270720      1.592131 -0.420888   1.786626 -1.520090

      OD  Proline
0      1.842721  1.010159
1      1.110317  0.962526
2      0.786369  1.391224
3      1.180741  2.328007
4      0.448336 -0.037767

```

```

..      ...      ...
173 -1.227742 -0.021890
174 -1.481267  0.009866
175 -1.481267  0.279786
176 -1.396759  0.295664
177 -1.424928 -0.593486

```

[178 rows x 14 columns]

```
[23]: df1.mean()
```

```

[23]: Wine                0.000000e+00
Alcohol                -9.181170e-16
Malic.acid             0.000000e+00
Ash                   -8.070947e-16
Acl                   -7.983626e-17
Mg                   -1.995907e-17
Phenols                3.991813e-17
Flavanoids            -3.592632e-16
Nonflavanoid.phenols  3.592632e-16
Proanth               -1.596725e-16
Color.int              1.995907e-17
Hue                   1.995907e-16
OD                    3.193450e-16
Proline               -7.983626e-17
dtype: float64

```

```

[24]: from sklearn.datasets import load_iris

iris = load_iris()

X = iris.data
Y = iris.target

iris_df = pd.DataFrame(X, columns=iris.feature_names)
iris_df

```

```

[24]:      sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0                5.1             3.5             1.4             0.2
1                4.9             3.0             1.4             0.2
2                4.7             3.2             1.3             0.2
3                4.6             3.1             1.5             0.2
4                5.0             3.6             1.4             0.2
..                ...             ...             ...             ...
145              6.7             3.0             5.2             2.3
146              6.3             2.5             5.0             1.9
147              6.5             3.0             5.2             2.0

```

| | | | | |
|-----|-----|-----|-----|-----|
| 148 | 6.2 | 3.4 | 5.4 | 2.3 |
| 149 | 5.9 | 3.0 | 5.1 | 1.8 |

[150 rows x 4 columns]

```
[25]: iris_df.std()
```

```
[25]: sepal length (cm)    0.828066
      sepal width (cm)     0.435866
      petal length (cm)    1.765298
      petal width (cm)     0.762238
      dtype: float64
```

```
[26]: iris_df.mean()
```

```
[26]: sepal length (cm)    5.843333
      sepal width (cm)     3.057333
      petal length (cm)    3.758000
      petal width (cm)     1.199333
      dtype: float64
```

```
[27]: df2 = (iris_df-iris_df.mean())/iris_df.std()
      df2

      df2.mean()
```

```
[27]: sepal length (cm)    -1.415442e-15
      sepal width (cm)     -1.652752e-15
      petal length (cm)    -1.442550e-15
      petal width (cm)     -5.543714e-16
      dtype: float64
```

```
[ ]:
```

```
[ ]:
```

Section 2: Data Mining Techniques

Run following algorithms on 2 real datasets and use appropriate evaluation measures to compute correctness of obtained patterns:

4 Practical Question 4

Ques4: Run Apriori algorithm to find frequent itemsets and association rules

1.1 Use minimum support as 50% and minimum confidence as 75%

1.2 Use minimum support as 60% and minimum confidence as 60 %

```
[28]: market_basket = pd.read_csv('Market_Basket_Optimisation.csv')
market_basket
```

```
[28]:
```

| | shrimp | almonds | avocado | vegetables mix | \ |
|------|----------------|-------------------|-------------|------------------|---|
| 0 | burgers | meatballs | eggs | NaN | |
| 1 | chutney | NaN | NaN | NaN | |
| 2 | turkey | avocado | NaN | NaN | |
| 3 | mineral water | milk | energy bar | whole wheat rice | |
| 4 | low fat yogurt | NaN | NaN | NaN | |
| ... | ... | ... | ... | ... | |
| 7495 | butter | light mayo | fresh bread | NaN | |
| 7496 | burgers | frozen vegetables | eggs | french fries | |
| 7497 | chicken | NaN | NaN | NaN | |
| 7498 | escalope | green tea | NaN | NaN | |
| 7499 | eggs | frozen smoothie | yogurt cake | low fat yogurt | |

| | green grapes | whole weat | flour yams | cottage cheese | energy drink | \ |
|------|--------------|------------|------------|----------------|--------------|---|
| 0 | NaN | NaN | NaN | NaN | NaN | |
| 1 | NaN | NaN | NaN | NaN | NaN | |
| 2 | NaN | NaN | NaN | NaN | NaN | |
| 3 | green tea | NaN | NaN | NaN | NaN | |
| 4 | NaN | NaN | NaN | NaN | NaN | |
| ... | ... | ... | ... | ... | ... | |
| 7495 | NaN | NaN | NaN | NaN | NaN | |
| 7496 | magazines | green tea | NaN | NaN | NaN | |
| 7497 | NaN | NaN | NaN | NaN | NaN | |
| 7498 | NaN | NaN | NaN | NaN | NaN | |
| 7499 | NaN | NaN | NaN | NaN | NaN | |

| | tomato juice | low fat yogurt | green tea | honey | salad | mineral water | salmon | \ |
|------|--------------|----------------|-----------|-------|-------|---------------|--------|---|
| 0 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 1 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 2 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 3 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 4 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 7495 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 7496 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 7497 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 7498 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |
| 7499 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | |

| | antioxydant juice | frozen smoothie | spinach | olive oil |
|---|-------------------|-----------------|---------|-----------|
| 0 | NaN | NaN | NaN | NaN |
| 1 | NaN | NaN | NaN | NaN |
| 2 | NaN | NaN | NaN | NaN |
| 3 | NaN | NaN | NaN | NaN |

| | | | | |
|------|-----|-----|-----|-----|
| 4 | NaN | NaN | NaN | NaN |
| ... | ... | ... | ... | ... |
| 7495 | NaN | NaN | NaN | NaN |
| 7496 | NaN | NaN | NaN | NaN |
| 7497 | NaN | NaN | NaN | NaN |
| 7498 | NaN | NaN | NaN | NaN |
| 7499 | NaN | NaN | NaN | NaN |

[7500 rows x 20 columns]

```
[29]: clean_data = []

for i in range(len(market_basket)):
    clean_data.append([x for x in list(map(str,market_basket.iloc[i].tolist()))
        ↪if x != 'nan'])
length = len(clean_data)
print(length)
```

7500

1.1 Use minimum support as 50% and minimum confidence as 75%

```
[30]: from efficient_apriori import apriori
itemsets, rules = apriori(clean_data, min_support=0.05, min_confidence=0.075,
    ↪verbosity=1)
```

Generating itemsets.

```
Counting itemsets of length 1.
Found 120 candidate itemsets of length 1.
Found 25 large itemsets of length 1.
Counting itemsets of length 2.
Found 300 candidate itemsets of length 2.
Found 3 large itemsets of length 2.
Counting itemsets of length 3.
Found 0 candidate itemsets of length 3.
Itemset generation terminated.
```

Generating rules from itemsets.

```
Generating rules of size 2.
Rule generation terminated.
```

```
[31]: itemsets
```

```
[31]: {1: {('burgers',): 654,
        ('eggs',): 1348,
        ('turkey',): 469,
        ('mineral water',): 1787,
```

```

('milk',): 972,
('whole wheat rice',): 439,
('green tea',): 990,
('low fat yogurt',): 573,
('french fries',): 1282,
('soup',): 379,
('frozen vegetables',): 715,
('spaghetti',): 1306,
('cookies',): 603,
('cooking oil',): 383,
('shrimp',): 535,
('chocolate',): 1229,
('chicken',): 450,
('tomatoes',): 513,
('pancakes',): 713,
('grated cheese',): 393,
('ground beef',): 737,
('frozen smoothie',): 474,
('escalope',): 595,
('cake',): 608,
('olive oil',): 493},
2: {('chocolate', 'mineral water'): 395,
    ('eggs', 'mineral water'): 382,
    ('mineral water', 'spaghetti'): 448}}

```

```

[32]: for item in sorted(rules, key=lambda item: (item.lift,item.conviction),
    ↪reverse=True):
    print(item)

```

```

{spaghetti} -> {mineral water} (conf: 0.343, supp: 0.060, lift: 1.440, conv:
1.159)
{mineral water} -> {spaghetti} (conf: 0.251, supp: 0.060, lift: 1.440, conv:
1.102)
{chocolate} -> {mineral water} (conf: 0.321, supp: 0.053, lift: 1.349, conv:
1.123)
{mineral water} -> {chocolate} (conf: 0.221, supp: 0.053, lift: 1.349, conv:
1.073)
{eggs} -> {mineral water} (conf: 0.283, supp: 0.051, lift: 1.189, conv: 1.063)
{mineral water} -> {eggs} (conf: 0.214, supp: 0.051, lift: 1.189, conv: 1.043)

```

1.2 Use minimum support as 60% and minimum confidence as 60 %

```

[33]: itemsets2, rules2 = apriori(clean_data, min_support=0.06, min_confidence=0.06,
    ↪verbosity=1)

```

Generating itemsets.

Counting itemsets of length 1.

Found 120 candidate itemsets of length 1.

```
Found 21 large itemsets of length 1.
Counting itemsets of length 2.
Found 210 candidate itemsets of length 2.
Itemset generation terminated.
```

```
Generating rules from itemsets.
Rule generation terminated.
```

```
[34]: itemsets2
```

```
[34]: {1: {('burgers',): 654,
        ('eggs',): 1348,
        ('turkey',): 469,
        ('mineral water',): 1787,
        ('milk',): 972,
        ('green tea',): 990,
        ('low fat yogurt',): 573,
        ('french fries',): 1282,
        ('frozen vegetables',): 715,
        ('spaghetti',): 1306,
        ('cookies',): 603,
        ('shrimp',): 535,
        ('chocolate',): 1229,
        ('chicken',): 450,
        ('tomatoes',): 513,
        ('pancakes',): 713,
        ('ground beef',): 737,
        ('frozen smoothie',): 474,
        ('escalope',): 595,
        ('cake',): 608,
        ('olive oil',): 493}}
```

```
[35]: for item in sorted(rules, key=lambda item: (item.lift,item.conviction),
        ↪reverse=True):
        print(item)
```

```
{spaghetti} -> {mineral water} (conf: 0.343, supp: 0.060, lift: 1.440, conv:
1.159)
{mineral water} -> {spaghetti} (conf: 0.251, supp: 0.060, lift: 1.440, conv:
1.102)
{chocolate} -> {mineral water} (conf: 0.321, supp: 0.053, lift: 1.349, conv:
1.123)
{mineral water} -> {chocolate} (conf: 0.221, supp: 0.053, lift: 1.349, conv:
1.073)
{eggs} -> {mineral water} (conf: 0.283, supp: 0.051, lift: 1.189, conv: 1.063)
{mineral water} -> {eggs} (conf: 0.214, supp: 0.051, lift: 1.189, conv: 1.043)
```

```
[ ]:
```

```
[ ]:
```

5 Practical Question 5

Ques5: Use Naive bayes, K-nearest, and Decision tree classification algorithms and build classifiers. Divide the data set into training and test set. Compare the accuracy of the different classifiers under the following situation:

```
[36]: from sklearn.datasets import load_iris
      from sklearn.model_selection import train_test_split ,
      ↪cross_val_score,KFold,StratifiedKFold
      from sklearn.naive_bayes import GaussianNB
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import accuracy_score
      from sklearn.preprocessing import StandardScaler

[37]: iris = load_iris()
      X = iris.data
      y = iris.target

      # Splitting the data into Training set = 75% & Test set = 25%
      X_train1, X_test1, y_train1, y_test1 = train_test_split(X, y, test_size=0.25,
      ↪random_state=0)

      # Splitting the data into Training set = 66.6% (2/3rd of total) & Test set =33.
      ↪3%
      X_train2, X_test2, y_train2, y_test2 = train_test_split(X, y, test_size=0.33,
      ↪random_state=0)

[38]: # initialize the models
      gnb = GaussianNB()
      knn = KNeighborsClassifier(n_neighbors=5)
      dt = DecisionTreeClassifier(max_depth=3)
```

5.1 a) Training set = 75% Test set = 25% b) Training set = 66.6% (2/3rd of total), Test set = 33.3

```
[39]: gnb.fit(X_train1, y_train1)
      y_pred1 = gnb.predict(X_test1)
      print("Naive Bayes accuracy for 75% training set: ", accuracy_score(y_test1,
      ↪y_pred1))

      knn.fit(X_train1, y_train1)
      y_pred2 = knn.predict(X_test1)
```

```

print("KNN accuracy for 75% training set: ", accuracy_score(y_test1, y_pred2))

dt.fit(X_train1, y_train1)
y_pred3 = dt.predict(X_test1)
print("Decision Tree accuracy for 75% training set: ", accuracy_score(y_test1,
    ↪y_pred3))

```

Naive Bayes accuracy for 75% training set: 1.0

KNN accuracy for 75% training set: 0.9736842105263158

Decision Tree accuracy for 75% training set: 0.9736842105263158

```

[40]: # 5.1 b) Training set = 66.6% (2/3rd of total), Test set =33.3%
gnb.fit(X_train2, y_train2)
y_pred4 = gnb.predict(X_test2)
print("Naive Bayes accuracy for 66.6% training set: ", accuracy_score(y_test2,
    ↪y_pred4))

knn.fit(X_train2, y_train2)
y_pred5 = knn.predict(X_test2)
print("KNN accuracy for 66.6% training set: ", accuracy_score(y_test2, y_pred5))

dt.fit(X_train2, y_train2)
y_pred6 = dt.predict(X_test2)
print("Decision Tree accuracy for 66.6% training set: ",
    ↪accuracy_score(y_test2, y_pred6))

```

Naive Bayes accuracy for 66.6% training set: 0.96

KNN accuracy for 66.6% training set: 0.98

Decision Tree accuracy for 66.6% training set: 0.94

5.2 Training set is chosen by i) hold out method ii) Random subsampling iii) Cross-Validation. Compare the accuracy of the classifiers obtained

```

[41]: # hold out method
X_train3 , X_test3 , y_train3 , y_test3 = train_test_split(X, y, test_size=0.
    ↪20, random_state=42)

gnb.fit(X_train3, y_train3)
y_pred7 = gnb.predict(X_test3)
print("Naive Bayes accuracy for hold out method: ", accuracy_score(y_test3,
    ↪y_pred7))

knn.fit(X_train3, y_train3)
y_pred8 = knn.predict(X_test3)
print("KNN accuracy for hold out method: ", accuracy_score(y_test3, y_pred8))

dt.fit(X_train3, y_train3)
y_pred9 = dt.predict(X_test3)

```

```
print("Decision Tree accuracy for hold out method: ", accuracy_score(y_test3, y_pred9))
```

Naive Bayes accuracy for hold out method: 1.0
KNN accuracy for hold out method: 1.0
Decision Tree accuracy for hold out method: 1.0

```
[42]: # Random subsampling method
sum_gnb = 0
sum_knn = 0
sum_dt = 0
for i in range(10):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=i)
    gnb.fit(X_train, y_train)
    y_pred = gnb.predict(X_test)
    sum_gnb += accuracy_score(y_test, y_pred)

    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    sum_knn += accuracy_score(y_test, y_pred)

    dt.fit(X_train, y_train)
    y_pred = dt.predict(X_test)
    sum_dt += accuracy_score(y_test, y_pred)

print("Naive Bayes accuracy for random subsampling method: ", sum_gnb/10)
print("KNN accuracy for random subsampling method: ", sum_knn/10)
print("Decision Tree accuracy for random subsampling method: ", sum_dt/10)
```

Naive Bayes accuracy for random subsampling method: 0.9473684210526315
KNN accuracy for random subsampling method: 0.9631578947368421
Decision Tree accuracy for random subsampling method: 0.944736842105263

```
[43]: # cross-validation
gnb_scores = cross_val_score(gnb, X, y, cv=10)
print("Naive Bayes accuracy for cross-validation: ", gnb_scores.mean())

knn_scores = cross_val_score(knn, X, y, cv=10)
print("KNN accuracy for cross-validation: ", knn_scores.mean())

dt_scores = cross_val_score(dt, X, y, cv=10)
print("Decision Tree accuracy for cross-validation: ", dt_scores.mean())
```

Naive Bayes accuracy for cross-validation: 0.9533333333333334
KNN accuracy for cross-validation: 0.9666666666666668
Decision Tree accuracy for cross-validation: 0.96

5.3 Data is scaled to standard formats:

```
[44]: scaler = StandardScaler()
      scaler.fit(X)
      X_scaled = scaler.transform(X)

      x_train4 , x_test4 , y_train4 , y_test4 = train_test_split(X_scaled, y,
      ↪test_size=0.20, random_state=42)

      gnb.fit(x_train4, y_train4)
      y_pred10 = gnb.predict(x_test4)
      print("Naive Bayes accuracy for scaled data: ", accuracy_score(y_test4,
      ↪y_pred10))

      knn.fit(x_train4, y_train4)
      y_pred11 = knn.predict(x_test4)
      print("KNN accuracy for scaled data: ", accuracy_score(y_test4, y_pred11))

      dt.fit(x_train4, y_train4)
      y_pred12 = dt.predict(x_test4)
      print("Decision Tree accuracy for scaled data: ", accuracy_score(y_test4,
      ↪y_pred12))
```

Naive Bayes accuracy for scaled data: 1.0

KNN accuracy for scaled data: 1.0

Decision Tree accuracy for scaled data: 1.0

[]:

[]:

6 Practical Question 6

Ques6: Use Simple Kmeans, DBScan, Hierarchical clustering algorithms for clustering. Compare the performance of clusters by changing the parameters involved in the algorithms.

```
[45]: import numpy as np
      import matplotlib.pyplot as plt
      from sklearn.datasets import make_blobs
      from sklearn.cluster import KMeans, DBSCAN, AgglomerativeClustering
      from sklearn.metrics import silhouette_score

      # Generating sample data
      X, _ = make_blobs(n_samples=300, centers=4, cluster_std=0.60, random_state=0)

      # Visualizing the sample data
```

```

plt.scatter(X[:, 0], X[:, 1], s=50)
plt.xlabel('X')
plt.ylabel('Y')
plt.title('Sample Data')
plt.show()

# Simple KMeans Clustering
def kmeans_clustering(X, n_clusters):
    kmeans = KMeans(n_clusters=n_clusters)
    kmeans.fit(X)
    labels = kmeans.labels_
    return labels

# DBSCAN Clustering
def dbscan_clustering(X, eps, min_samples):
    dbscan = DBSCAN(eps=eps, min_samples=min_samples)
    labels = dbscan.fit_predict(X)
    return labels

# Hierarchical Clustering
def hierarchical_clustering(X, n_clusters):
    agg_clustering = AgglomerativeClustering(n_clusters=n_clusters)
    labels = agg_clustering.fit_predict(X)
    return labels

# Evaluating clustering performance using silhouette score
def evaluate_clustering(X, labels):
    if len(np.unique(labels)) > 1:
        silhouette_avg = silhouette_score(X, labels)
        return silhouette_avg
    else:
        return -1 # Return a placeholder value when only noise points are
↳ present

# Testing Simple KMeans with different number of clusters
print("Simple KMeans:")
for n_clusters in range(2, 6):
    labels = kmeans_clustering(X, n_clusters)
    silhouette_avg = evaluate_clustering(X, labels)
    print(f"Number of clusters: {n_clusters}, Silhouette Score:↳
↳ {silhouette_avg}")

# Testing DBSCAN with different values of epsilon and min_samples
print("\nDBSCAN:")
eps_values = [0.3, 0.5, 0.7, 1.0]
min_samples_values = [5, 10, 15, 20]
for eps in eps_values:

```




```

for min_samples in min_samples_values:
    labels = dbscan_clustering(X, eps, min_samples)
    silhouette_avg = evaluate_clustering(X, labels)
    if silhouette_avg != -1:
        print(f"EPS: {eps}, Min Samples: {min_samples}, Silhouette Score:␣
↪{silhouette_avg}")
    else:
        print(f"EPS: {eps}, Min Samples: {min_samples}, No valid clusters␣
↪found.")

# Testing Hierarchical Clustering with different number of clusters
print("\nHierarchical Clustering:")
for n_clusters in range(2, 6):
    labels = hierarchical_clustering(X, n_clusters)
    silhouette_avg = evaluate_clustering(X, labels)
    print(f"Number of clusters: {n_clusters}, Silhouette Score:␣
↪{silhouette_avg}")

```



Simple KMeans:

Number of clusters: 2, Silhouette Score: 0.5426422297358302

Number of clusters: 3, Silhouette Score: 0.5890390393551768
Number of clusters: 4, Silhouette Score: 0.6819938690643478
Number of clusters: 5, Silhouette Score: 0.5923875148758644

DBSCAN:

EPS: 0.3, Min Samples: 5, Silhouette Score: -0.02553097772433596
EPS: 0.3, Min Samples: 10, Silhouette Score: -0.25081172884728203
EPS: 0.3, Min Samples: 15, No valid clusters found.
EPS: 0.3, Min Samples: 20, No valid clusters found.
EPS: 0.5, Min Samples: 5, Silhouette Score: 0.6303800996842714
EPS: 0.5, Min Samples: 10, Silhouette Score: 0.5220954071399261
EPS: 0.5, Min Samples: 15, Silhouette Score: 0.3871688267990456
EPS: 0.5, Min Samples: 20, Silhouette Score: 0.2097450211809529
EPS: 0.7, Min Samples: 5, Silhouette Score: 0.559707233404896
EPS: 0.7, Min Samples: 10, Silhouette Score: 0.6569398552813946
EPS: 0.7, Min Samples: 15, Silhouette Score: 0.6198214828463691
EPS: 0.7, Min Samples: 20, Silhouette Score: 0.6045937193339364
EPS: 1.0, Min Samples: 5, Silhouette Score: 0.46285745923867483
EPS: 1.0, Min Samples: 10, Silhouette Score: 0.58977292182092
EPS: 1.0, Min Samples: 15, Silhouette Score: 0.58977292182092
EPS: 1.0, Min Samples: 20, Silhouette Score: 0.58977292182092

Hierarchical Clustering:

Number of clusters: 2, Silhouette Score: 0.54731479631826
Number of clusters: 3, Silhouette Score: 0.58977292182092
Number of clusters: 4, Silhouette Score: 0.6819938690643478
Number of clusters: 5, Silhouette Score: 0.5875473435823221

[]:

[]: